LEVERAGING CONVOLUTIONAL NEURAL NETWORKS FOR AUTOMATED PNEUMONIA DETECTION IN CHEST X-RAY IMAGES

Submitted in partial fulfillment of the

Requirement for the Master's degree

in

Information Technology

By

Rishani V

23/PCSA/101

Dec 2024 – Apr 2025

Stella Maris College (Autonomous)

17, Cathedral Road,

Chennai-600086.

STELLA MARIS COLLEGE (Autonomous)

17, Cathedral Road,

Chennai-600086.

Master of Science (Information Technology)

(Affiliated to University of Madras)

BONAFIDE CERTIFICATE

This is to certify that this is a bonafide record of the project done by

Rishani V

on

Predicting Student Performance through Classification

at

Stella Maris College (Autonomous)

Dec 2024 – Apr 2025

This project was done by her in partial fulfillment of the requirements for the Master's degree in Information Technology

Head of Department **Internal Examiner** **External Examiner**

Abstract

Pneumonia is a lung infection caused by bacteria, viruses, or fungi, leading to inflammation and fluid buildup in the lungs, making it hard to breathe. It's usually diagnosed using chest X-rays, but this process can take time and depends on the radiologist's expertise. This project focuses on creating a user-friendly application that uses Convolutional Neural Networks (CNNs) to automatically detect pneumonia. The CNN model, trained on a large dataset named chest X-ray images from Kaggle, can quickly and accurately differentiate between healthy individuals and those with pneumonia. Users can upload a chest X-ray image through the application and instantly receive a diagnosis. This tool provides a faster, reliable, and accessible way to support healthcare professionals in detecting pneumonia efficiently.

Keywords: *Pneumonia detection, Deep Learning, CNN, Chest X-ray images, Automated diagnosis.*

Table of Contents

Chapter 1: Introduction ... 4

1.1 Pneumonia Detection Using Deep Learning ... 4

1.2 Deep Learning in Medical Imaging ... 4

1.3 VGG16 for Pneumonia Detection ... 4

1.4 Dataset and Implementation ... 5

1.5 AI-Powered Healthcare Solutions ... 5

Chapter 2: Review of Past Work and Problem Formulation ... 6

2.1 Review of Previous Studies on Deep Learning for Pneumonia Detection ... 6

2.2 Study on CNN-Based Pneumonia Detection ... 6

2.3 Comparative Analysis of VGG16 and VGG19 ... 6

2.4 Dataset and Evaluation Metrics ... 7

2.4 Addressing Limitations of VGG16 ... 7

2.5 Implementation of Pre-Trained VGG16 Model ... 7

2.6 Practical Considerations for Pneumonia Detection ... 8

Chapter 3: Methodology ... 9

3.1 Data Preparation and Preprocessing ... 9

3.2 Model Training Using VGG16 ... 9

3.3 Web Application Development ... 10

Chapter 4: Results and Discussion/Analysis ... 14

Chapter 5: Conclusions and Scope for Future Work ... 15

Chapter 1: Introduction

1.1 Pneumonia Detection Using Deep Learning

Pneumonia is a severe lung infection that occurs in millions of people globally, causing serious health complications and fatalities, especially in children, the elderly, and immunocompromised individuals. Traditional diagnosis involves radiologists reading chest X-ray images, which is time-consuming and prone to human error. With the growing amount of medical imaging data, there is an urgent need for accurate, efficient, and automated diagnostic systems to aid doctors in early pneumonia detection.

1.2 Deep Learning in Medical Imaging

Deep learning has transformed medical imaging by making it possible to automatically diagnose diseases at high accuracy. Deep learning methods, Convolutional Neural Networks (CNNs), have also proven to be exceptionally effective for image classification. CNNs can learn from images and extract the critical visual features automatically without the need for heavy handcrafted feature engineering, hence best suited for medical image analysis. These networks have several layers, such as convolutional layers, pooling layers, and fully connected layers, which enable them to learn hierarchical patterns from images. CNNs have been applied extensively in medical conditions, such as pneumonia detection, cancer diagnosis, and retinal disease classification.

1.3 VGG16 for Pneumonia Detection

In this project, we utilize the VGG16 model, a pre-trained deep neural network with the characteristic of having a deep structure and hierarchical feature extraction. VGG16 is specifically suited for image classification, which makes it a good candidate for detecting pneumonia. The model has 16 layers, mostly convolutional and fully connected layers, and uses pre-trained on Chest X-ray images dataset from kaggle. Through transfer learning, VGG16 detects pneumonia-specific abnormalities with great accuracy while minimizing computational complexity and enhancing performance.

1.4 Dataset and Implementation

It is trained on Kaggle's Chest X-Ray Image Dataset, a set of chest X-ray images labeled as either normal or pneumonia-infected. Transfer learning employing VGG16 is used to take advantage of previously acquired knowledge to improve accuracy and efficiency. Automated classification goes a long way in preventing diagnostic failure and is an invaluable decision-support assistant for radiologists.

1.5 AI-Powered Healthcare Solutions

By incorporating deep learning in medical diagnostics, this research improves AI-based healthcare solutions. Employing CNNs, especially VGG16, demonstrates the prospect of computer-aided diagnosis of diseases to reduce radiologists' workload and facilitate quicker and more accurate classification of pneumonia. This research also emphasizes the revolutionary impact of artificial intelligence in contemporary medicine, with the potential to create more credible and effective diagnostic systems.

Chapter 2: Review of Past Work and Problem Formulation

2.1 Review of Previous Studies on Deep Learning for Pneumonia Detection

Several research studies have explored the application of deep learning, particularly Convolutional Neural Networks (CNNs), for pneumonia detection through chest X-ray images. These studies emphasize the efficacy of CNNs in medical imaging and highlight how deep learning models can automate pneumonia detection with high accuracy. Traditional pneumonia diagnosis using X-ray scans relies on radiologists, who may face challenges in distinguishing pneumonia from similar lung conditions. The increasing adoption of AI-based solutions aims to enhance diagnostic accuracy and reduce human dependency.

2.2 Study on CNN-Based Pneumonia Detection

One notable study, "A Study on Pneumonia Detection using CNN and a Comparative Analysis of VGG16 and VGG19," provides an in-depth evaluation of VGG16 and VGG19 architectures in classifying pneumonia. The study investigates how these CNN architectures extract meaningful features from X-ray images to improve classification accuracy. It underscores the importance of automated feature extraction in medical diagnostics, reducing the need for manual interpretation and minimizing diagnostic errors.

2.3 Comparative Analysis of VGG16 and VGG19

A key aspect of this research is the comparison between VGG16 and VGG19. While both models demonstrate strong performance in pneumonia detection, VGG19, with its deeper architecture, can capture more intricate image features. However, the study notes that deeper networks do not always result in significantly higher accuracy and may introduce challenges such as overfitting, increased computational demands, and longer training times. The research, which used the Kaggle Chest X-ray Image Dataset, found that although VGG19 achieved slightly better accuracy, VGG16 remains a practical choice due to its balance between computational efficiency and performance.

2.4 Dataset and Evaluation Metrics

The Kaggle Chest X-ray Image Dataset was used in this study, comprising labelled images classified as either normal or pneumonia-affected cases. Both VGG16 and VGG19 were trained on this dataset and evaluated using metrics such as accuracy, sensitivity, and specificity. The results demonstrated that although VGG19 performed slightly better in terms of accuracy, VGG16 remains an effective model due to its lower computational cost and faster training time.

2.4 Addressing Limitations of VGG16

While VGG16 is a widely used architecture for pneumonia detection, it has some limitations, particularly its sensitivity to image quality variations, which can impact classification performance. VGG19, with its deeper structure, can better handle these variations but at the cost of increased computational requirements and longer training times. To enhance VGG16's performance, preprocessing methods such as data rescaling using ImageDataGenerator from tensorflow.keras.preprocessing.image were applied. This step normalizes pixel values, stabilizes training, improves feature extraction, and enhances generalization, making the model more robust to variations in chest X-ray images.

2.5 Implementation of Pre-Trained VGG16 Model

To streamline implementation,VGG16 is imported directly from tensorflow.keras.applications.vgg16 instead of being manually defined layer by layer. This approach offers several advantages:

- Faster Implementation: Eliminates the need for manual layer specification, reducing development time.
- Reduced Errors: Ensures the correct architecture is used without human errors.

- Utilization of Pre-Trained Weights: Facilitates transfer learning using ImageNet weights, enabling faster convergence and improved feature learning.

2.6 Practical Considerations for Pneumonia Detection

Despite the advantages of deep learning models in medical imaging, real-world deployment requires careful consideration of computational efficiency, model interpretability, and clinical validation. While VGG19 offers slight improvements in accuracy, its higher computational cost makes it less practical for real-time applications. In contrast, VGG16, with proper preprocessing techniques and transfer learning, provides a more optimized and scalable solution for pneumonia detection. The integration of AI in medical imaging continues to evolve, with ongoing research focusing on improving model robustness, interpretability, and real-world applicability. Future work may explore hybrid deep learning architectures, improved data augmentation techniques, and real-time implementation strategies for automated pneumonia screening. Additionally, cloud-based and mobile AI solutions can further enhance accessibility, particularly in remote or resource-limited healthcare settings.

Chapter 3: Methodology

This project is designed to detect pneumonia from chest X-ray images using deep learning with VGG16 and a Flask-based web application. The system classifies X-ray images into two categories: Healthy or Pneumonic.

3.1 Data Preparation and Preprocessing

The dataset is structured into three directories: train, test, and validation. The ImageDataGenerator function is used to preprocess the images by rescaling pixel values to a range of 0-1. Additionally, for training data, augmentation techniques such as horizontal flipping, zooming, and shearing are applied to enhance model generalization. The validation and test datasets are only rescaled to ensure accurate evaluation. Images are loaded efficiently allowing batch processing.

```python
train_gen = ImageDataGenerator(rescale=1/255.0,
                               horizontal_flip=True,
                               zoom_range=0.2,
                               shear_range=0.2) # making the data loader for training data
test_gen = ImageDataGenerator(rescale=1/255.0) # making the data loader for validation data

train_data_gen = train_gen.flow_from_directory(train_dir,
                                               target_shape,
                                               batch_size=16,
                                               class_mode='categorical') # function to make iterable object for training
test_data_gen = train_gen.flow_from_directory(test_dir,
                                              target_shape,
                                              batch_size=16,
                                              class_mode='categorical') # function to make iterable object for training
```

Fig 3.3.1.1 Data Preprocessing

3.2 Model Training Using VGG16

A pre-trained VGG16 model is used with transfer learning to extract meaningful features from X-ray images. The model is loaded without the fully connected layers and a new classification head is added. The final architecture includes a Flatten layer, followed by a Dense layer with Softmax activation for binary classification. The model is compiled using the Adam optimizer, with categorical cross-entropy as the loss function. Training is performed using the augmented dataset, and accuracy/loss graphs are plotted to track performance. Once trained, the model is saved as model.h5 for later deployment.

```python
# loading the VGG16 model
vgg = VGG16(weights='imagenet',include_top=False,input_shape=(224,224,3))
for layer in vgg.layers:
    layer.trainable = False #making all the layers non-trainable

x = Flatten()(vgg.output) #flattening out the last layer
predictions = Dense(2,activation='softmax')(x) #Dense layer to predict wether their is pneumonia or not
model = Model(inputs=vgg.input, outputs=predictions)
model.summary()
```

Fig 3.3.2.2 Model Training

3.3 Web Application Development

A Flask-based web application is developed to provide an easy-to-use interface for pneumonia detection. Users can register and log in. The web application allows users to upload chest X-ray images, which are processed by the trained model to predict whether the image is Healthy or Pneumonic. The results are displayed on the home page, with a clean and responsive design using Bootstrap. The system includes multiple web pages such as index.html (homepage), login.html (user authentication), register.html (new user registration), and home.html (image upload and result display).

index.html

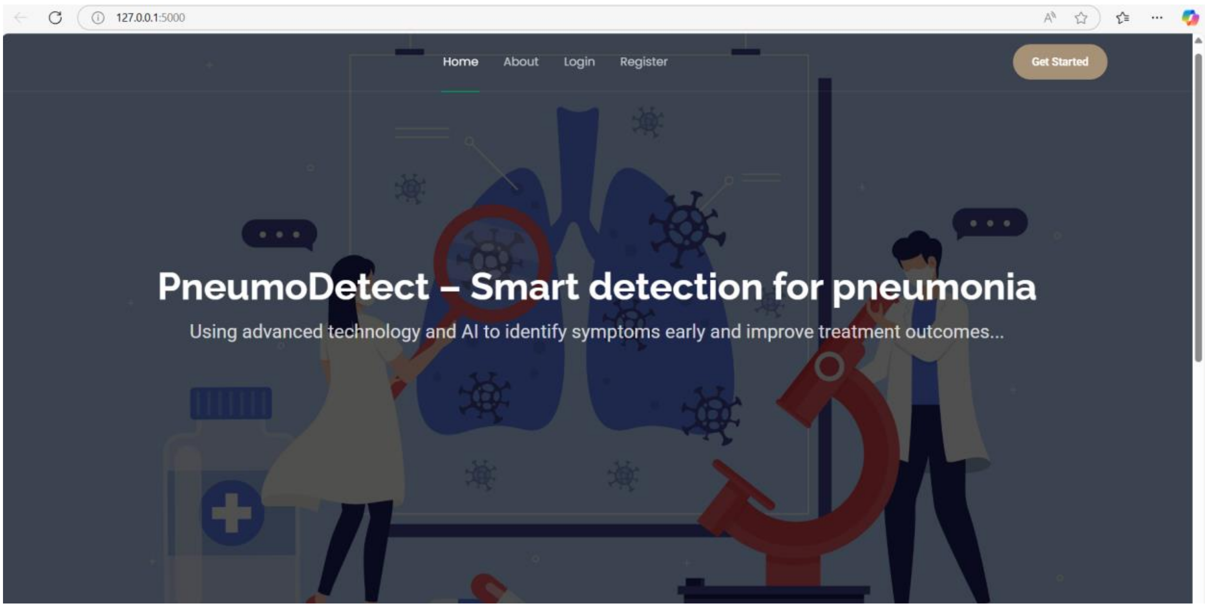

Fig 3.3.1 Home page

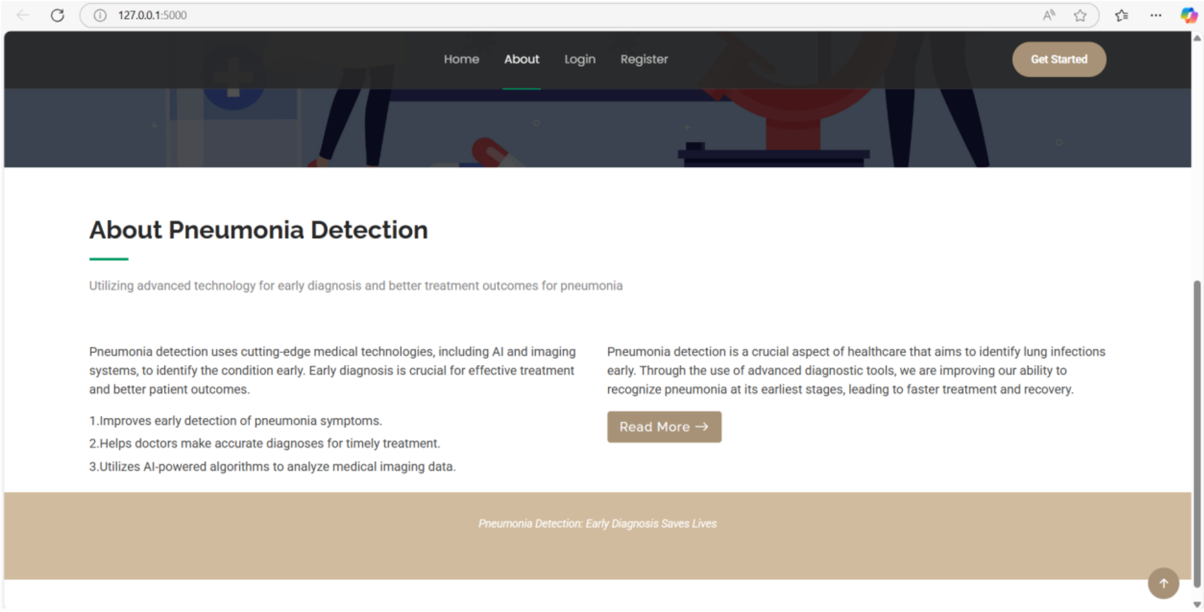

Fig 3.3.2 Home page about details

login.html

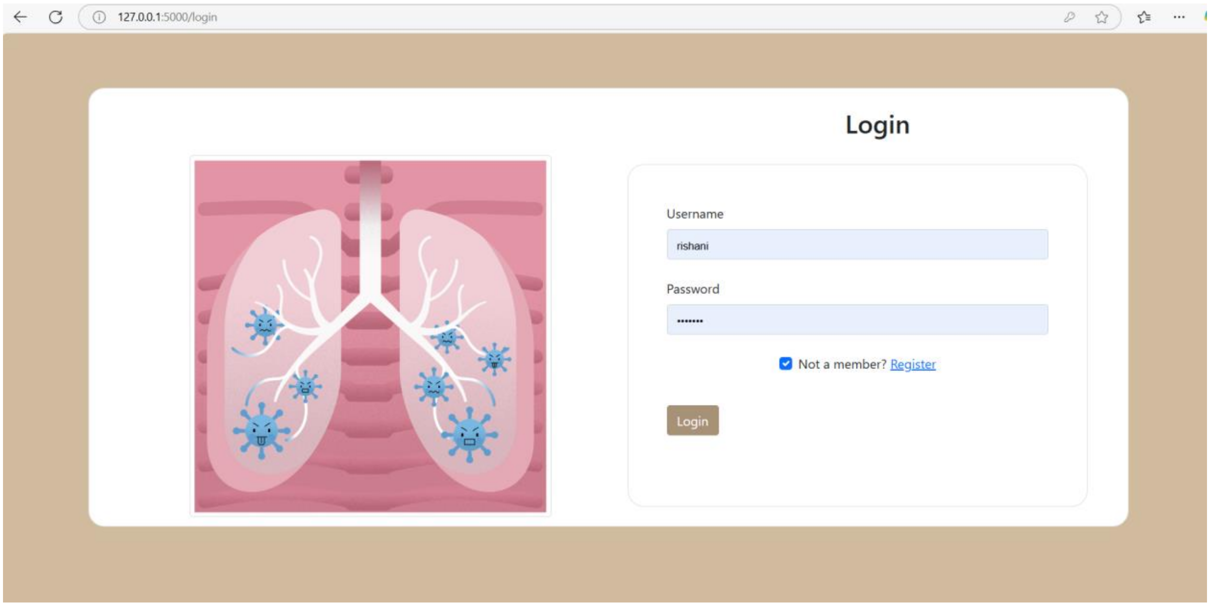

Fig 3.3.3 Login page

register.html

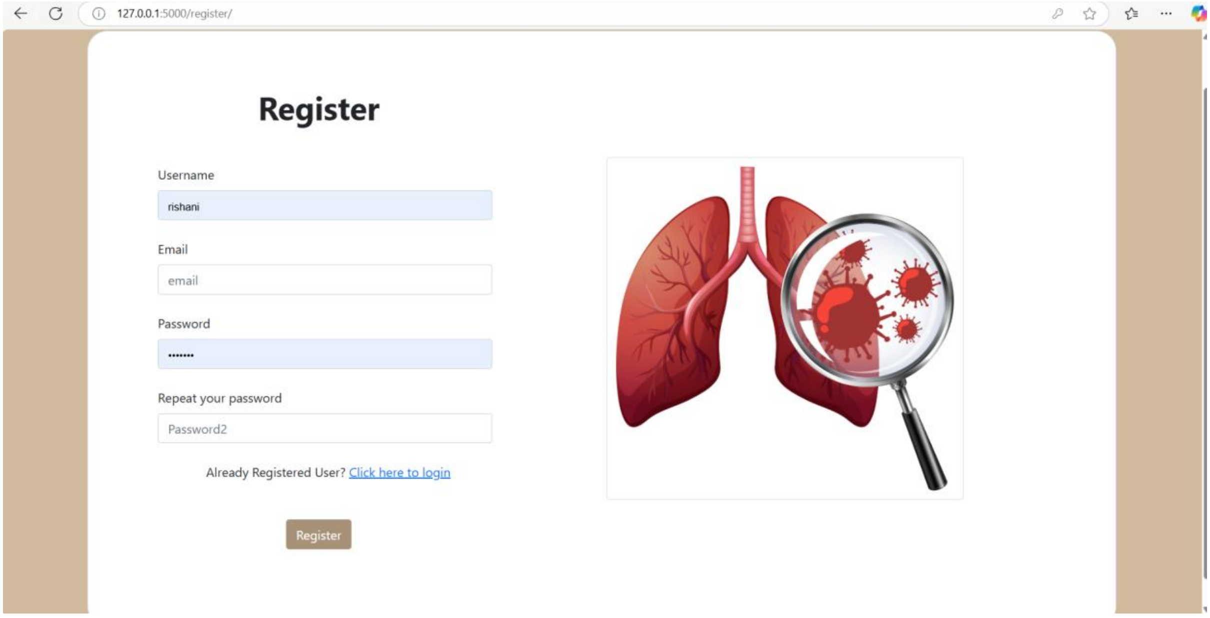

Fig 3.3.4 Register page

home.html

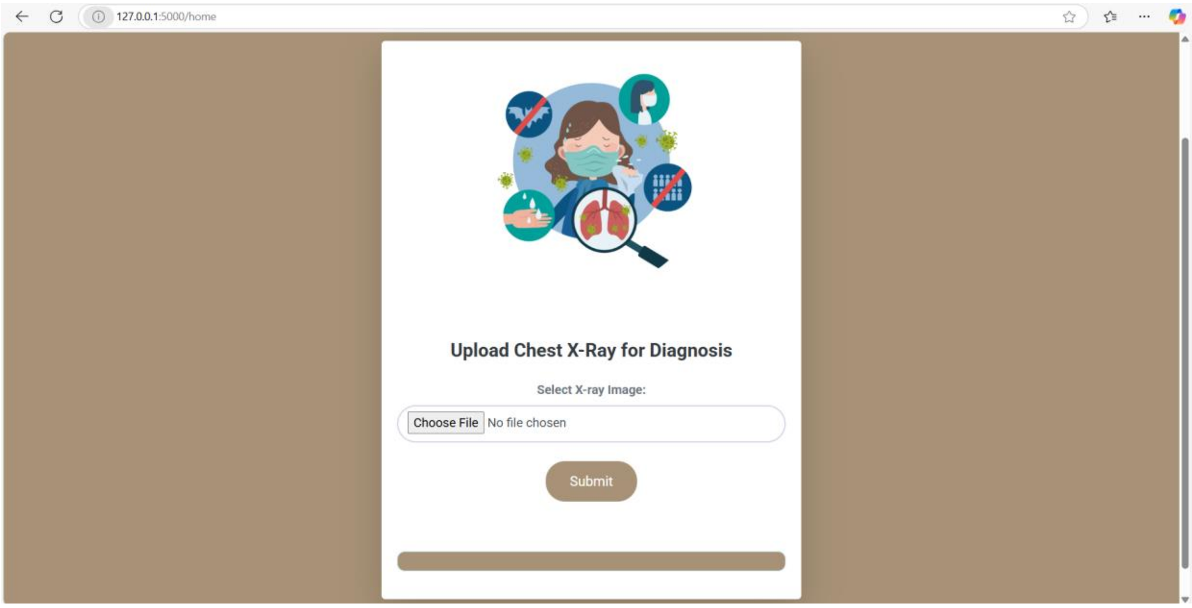

Fig 3.3.5 Image Upload page for Pneumonia detection

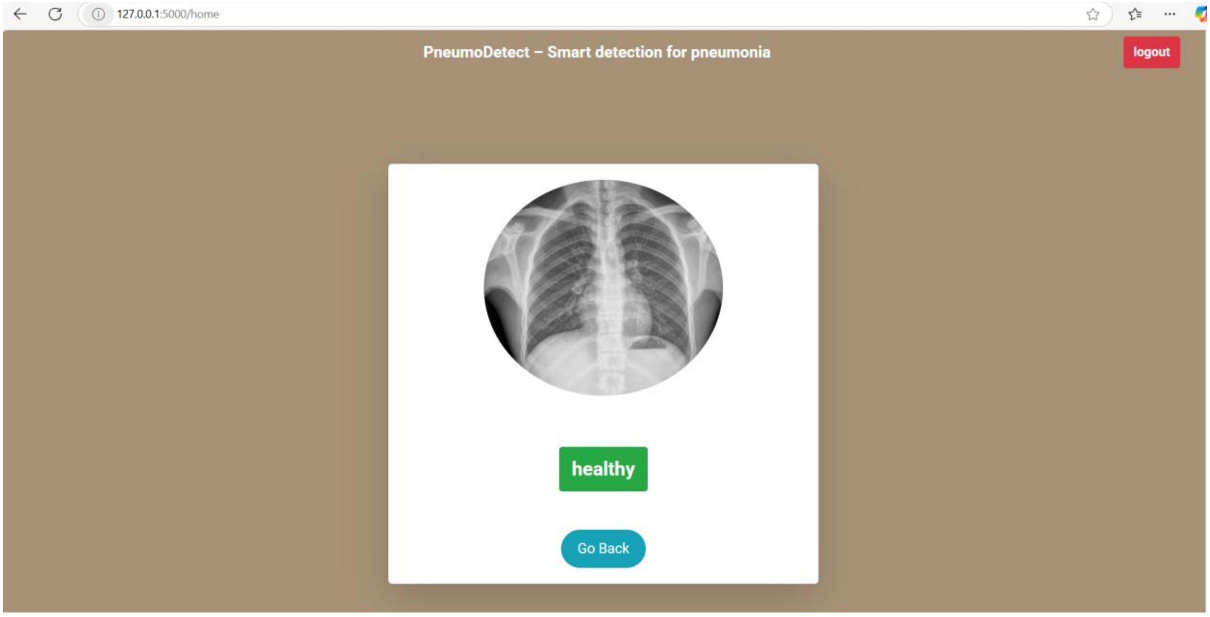

Fig 3.3.6 Healthy Chest X-ray

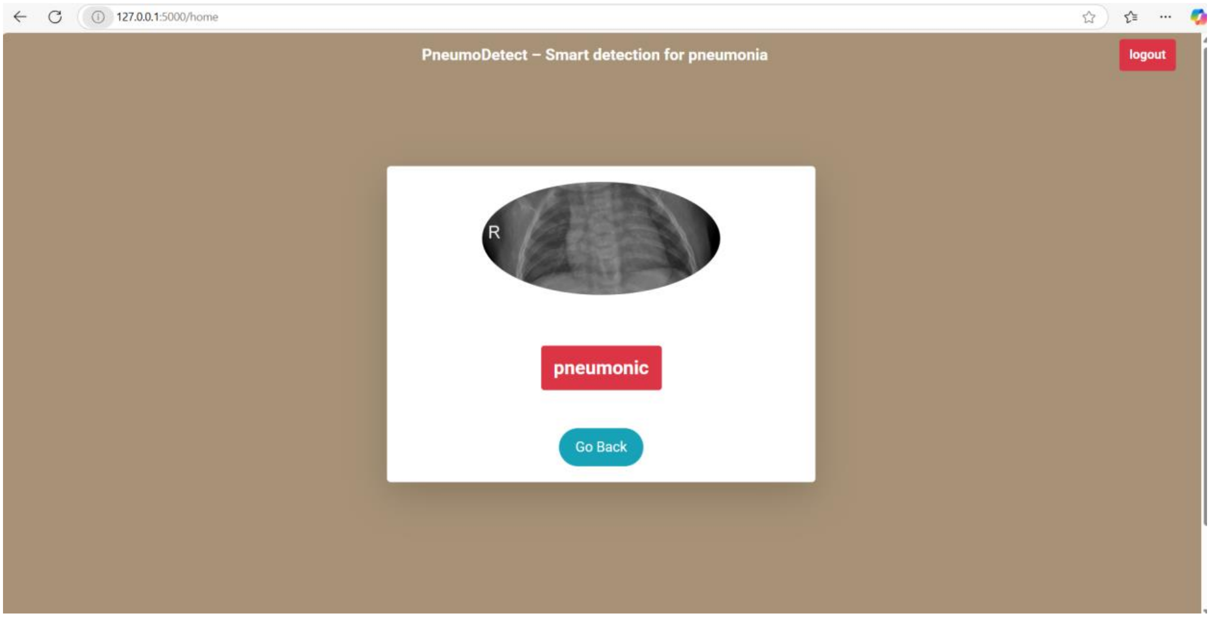

Fig 3.3.7 Pneumonia Chest X-ray

Chapter 4: Results and Discussion/Analysis

The pneumonia detection system was evaluated based on its accuracy, efficiency, and usability. The model was trained using VGG16 with transfer learning, and ImageDataGenerator was used to improve performance by applying techniques like rescaling, flipping, and zooming. The training process showed steady improvement in accuracy over multiple epochs, with the final model achieving good classification results on the test dataset.

During testing, the model correctly classified most chest X-ray images as Healthy or Pneumonic. However, some errors occurred due to poor image quality or similarities between pneumonia and other lung conditions. Despite this, the evaluation metrics, including precision, recall, and F1-score, showed that the model effectively detected pneumonia while keeping false predictions low.

The Flask-based web application allowed users to easily upload X-ray images and get instant results. The web interface was tested for usability, and predictions were displayed in a clear and understandable format, making it accessible for both medical professionals and general users. While the system is effective, real-world medical deployment would require further testing with larger and more diverse datasets to ensure consistent accuracy across different patient cases.

Overall, the system performed well in detecting pneumonia, offering a fast and simple solution for early diagnosis. Future improvements, such as expanding the dataset, enhancing security, and optimizing the model for mobile devices, can further improve accuracy and usability, making the system a more reliable tool for pneumonia detection in healthcare settings.

Chapter 5: Conclusions and Scope for Future Work

This project successfully uses deep learning and a web application to detect pneumonia from chest X-ray images. The system, built with VGG16 and transfer learning, can classify images as Healthy or Pneumonic and provides a simple, user-friendly interface for diagnosis. The use of ImageDataGenerator helps improve accuracy by making the model better at handling different types of images.

In the future, the system can be improved by using a larger and more diverse dataset, connecting it with hospital systems for real-time use, and making it work on mobile devices for easier access. Better security and privacy features can be added to protect patient data. Adding explainable AI will help doctors understand why the model makes certain predictions. These improvements will make the system more accurate, accessible, and useful for healthcare professionals and patients.